

Jagadish B. Kotra

Researcher, Member of Technical Staff, AMD Research, Austin, Texas

Website: <https://jbk5155.github.io/>

Email: jagadishkotra@gmail.com

Mobile: (813)-468-8220

RESEARCH INTERESTS

Computer Architecture, Operating Systems, Hardware-OS co-design, Heterogeneous CPU/GPU Systems.

EDUCATION

Doctor of Philosophy, The Pennsylvania State University

2010 - 2017

Advisor: **Dr. Mahmut T. Kandemir**.

PhD Dissertation Topic: **Hardware-Software co-design for optimizing memory-hierarchy in large and many-core systems.**

Bachelor of Technology, Acharya Nagarjuna University

2002 - 2006

Major: Electronics and Communications Engineering

PUBLICATIONS

- "Improving the Utilization of Micro-operations cache in x86 Processors". **Jagadish Kotra**, John Kalamatianos [MICRO 2020]
- "DSM: A Case for Hardware-Assisted Merging of DRAM Rows with Same Content". Armin Vakil, Mahmut Kandemir, **Jagadish Kotra** [SIGMETRICS 2020]
- "Centaur: A Novel Architecture for Reliable, Low-Wear, High-Density 3D NAND Storage". Chun-yi Liu, **Jagadish Kotra**, Myoungsoo Jung, Mahmut Kandemir. [SIGMETRICS 2020]
- "PreFAM: Understanding the Impact of Prefetching in Fabric-Attached Memory Architectures". Vamsee R.K, **Jagadish Kotra**, Clayton H., Hammond S.D, Amro Awad. [MemSys 2020]
- "Optimization of Inter-Cache Traffic Entanglement in Tagless Caches with Tiling Opportunities". S.R. Swamy, Sumitha G., Hariram G., **Jagadish Kotra**, Madhu M., Jack S., Mahmut K., Vijay N. [CASES 2020]
- "CASH: Compiler ASsisted Hardware Design for improving DRAM energy efficiency in CPU-based Inference Systems". Anup Sarma, Huaipann Jiang, Ashutosh Pattnaik, **Jagadish Kotra**, Mahmut Kandemir, Chita Das [MemSys 2019]
- "SOML Read: Rethinking the read operation granularity of 3D NAND SSDs". Chun-yi Liu, **Jagadish Kotra**, Myoungsoo Jung, Mahmut Kandemir, Chita R. Das. [ASPLOS 2019]
- "CHAMELEON: A co-design based dynamically reconfigurable heterogeneous memory system". **Jagadish Kotra**, Haibo Zhang, Alaa R. Alameldeen, Chris Wilkerson, Mahmut Kandemir. [MICRO 2018]
- "MDACache: Caching for Multi-Dimensional-Access Memories". Sumitha George, Minli Liao, Huaipan Jiang, **Jagadish Kotra**, Mahmut Kandemir, John Sampson, Vijaykrishnan Narayanan. [MICRO 2018]
- "PEN: A Design of Partial-Erase for 3D NAND-based High Capacity SSDs". Chun-Yi Liu, **Jagadish Kotra**, Myoungsoo Jung, Mahmut Kandemir. [FAST 2018]
- "Enhancing Computation-to-Core Assignment with Physical Location Information". Orhan Kislal, **Jagadish Kotra**, Xulong Tang, Mahmut Kandemir, Myoungsoo Jung, Mustafa Karakoy. [PLDI 2018]
- "A Learning-guided Hierarchical Approach for Biomedical Image Segmentation". Huaipan Jiang, Anup Sarma, Jihyun Ryoo, **Jagadish Kotra**, Meena A., Chita Das, Mahmut Kandemir. [SoCC 2018]
- "Hardware-software co-design to mitigate DRAM refresh overheads: A case for DRAM refresh-aware process scheduling". **Jagadish Kotra**, Narges S., Zeshan Chishti, Mahmut Kandemir. [ASPLOS 2017]

- “Congestion Aware Memory Management on NUMA platforms: A VMware ESXi case study”. **Jagadish Kotra**, Seongbeom Kim, Kamesh Madduri, Mahmut Kandemir. [IISWC 2017]
- “Quantifying the Potential Benefits of Near-Data Computing in Manycores”. **Jagadish Kotra**, Diana Guttman, Nachiappan C, Mahmut Kandemir, Chita Das. [MASCOTS 2017]
- “Location-Aware Computation Mapping for Manycores”. Orhan Kislal, **Jagadish Kotra**, Xulong Tang, Mahmut Kandemir, Myoungsoo Jung, Mustafa Karakoy. [PACT 2017]
- “Re-NUCA: A practical NUCA architecture for Re-RAM based last-level caches”. **Jagadish Kotra**, Mohammed Arjomand, Diana Guttman, Mahmut Kandemir, Chita Das. [IPDPS 2016]
- “Improving Bank-Level Parallelism for Irregular Applications”. Xulong Tang, Mahmut Kandemir, Praveen Yedlapalli, **Jagadish Kotra**. [MICRO 2016]
Best paper award nominee
- “Cache-Aware Approximate Computing for Decision Tree Learning”. Orhan Kislal, Mahmut Kandemir, **Jagadish Kotra**. [IPDPS-Parlearning 2016]
- “Thermal-aware Application Scheduling on Device-heterogeneous Embedded Architectures”. Karthik Swaminathan, **Jagadish Kotra**, Mahmut Kandemir and Vijaykrishnan Narayanan. [VLSID 2015]
- “Network Footprint Reduction through Data Access and Computation Placement in NoC-Based Many cores”. Jun Liu, **Jagadish Kotra**, Wei Ding, Mahmut Kandemir. [DAC 2015]
- “Phase Detection with Hidden Markov Models for DVFS on Many-Core Processors”. Joshua Booth, **Jagadish Kotra**, Hui Zhao, Mahmut Kandemir, Padma Raghavan. [ICDCS 2015]
- “Meeting Midway: Improving DRAM Performance and Off-Chip Latencies with Memory-Side Prefetching”. Praveen Yedlapalli, **Jagadish Kotra**, Emre Kultursay, Mahmut Kandemir, Anand Sivasubramaniam, Chita R. Das. [PACT 2013]

CONFERENCE TALKS

- “Congestion-aware Memory Management: A VMWare ESXi case study”. Seattle, USA. IISWC-2017.
- “Hardware-Software co-design: A case for refresh-aware process scheduling”. Xian, China. ASPLOS-2017.
- “Quantifying the potential benefits of Near-data Computing in Manycores”. Banff, Quebec, Canada. MASCOTS-2017.
- “Re-NUCA: A NUCa architecture for Re-RAM based last-level caches”. Chicago, USA. IPDPS-2016.
- “Cache-Aware Approximate Computing for Decision Tree Learning”. Chicago, USA. IPDPS-2016.
- “Network Footprint Reduction through Data Access and Computation Placement in NoC-Based Many cores”. San Francisco, USA. DAC-2016.

U.S PATENTS

- “A hardware-assisted DRAM row merging mechanism for energy-efficiency”. **Jagadish Kotra**. (About to be filed in USPTO, AMD)
- “A method and apparatus for reducing average latency of long latency load instructions”. **Jagadish Kotra**, John Kalamatianos. (About to be filed in USPTO, AMD).
- “A case for atomics arbitration”. Sergey Blagodurov, John Alsop, **Jagadish Kotra**, Marko Scrbak. (About to be filed in USPTO, AMD)
- “GPU Reach Optimizations”. **Jagadish Kotra**, Michael Lebeane (About to be filed in USPTO, AMD)
- “Method and Apparatus for Speculative Data Promotion from the Cache to the Physical Register File”. **Jagadish Kotra**, John Kalamatianos. (About to be filed in USPTO, AMD)
- 170499-US-NP. “A method and apparatus for Optimizing Micro-op Cache”. **Jagadish Kotra**, John Kalamatianos (filed in USPTO, on behalf of AMD).
- “Hardware-software collaborative address mapping scheme for efficient processing-in-memory systems”. Mahzabeen I., Shaizeen A., Nuwan J., **Jagadish Kotra**. (About to be filed in USPTO, AMD)

- 180128-US-NP. "A Method and Apparatus for temperature-gradient aware data-placement for 3D stacked DRAMs in GPUs". **Jagadish Kotra**, Karthik R., J. Greathouse. (filed in USPTO, on behalf of AMD).
- 180399-US-NP. "Method and apparatus for improving the utilization of micro-op caches via compaction". **Jagadish Kotra**, John Kalamatianos. (filed in USPTO, on behalf of AMD)
- 180243-US-NP. "A Method for a Generative Adversarial Network Resource Scheduler." Sergey Blagodurov, Abhinav Vishnu, Thaleia dimitra Doudali, **Jagadish Kotra**. (filed in USPTO, on behalf of AMD)
- "Method and Apparatus for Speculative Data Promotion from the Cache to the Physical Register File". **Jagadish Kotra**, John Kalamatianos. (About to be filed in USPTO, AMD)
- "Methods for Configuring Span of Control Under Varying Temperature". Tony Gutierrez, Yasuko Eckert, Sergey Blagodurov, **Jagadish Kotra**. (About to be filed in USPTO, AMD)
- "Micro-operations cache Allocation Filter". **Jagadish Kotra**, Marko Scrback, Mahzabeen Islam, John Kalamatianos (filed in USPTO, on behalf of AMD)
- "Mechanisms for Temporal Link Encoding". Onur Kayiran, Steve Raasch, Sergey Blagodurov, **Jagadish Kotra**. (filed in USPTO, on behalf of AMD)
- US 2017/0371777 A1. "Memory Congestion Aware NUMA Management". **Jagadish Kotra**, S. Kim, Fei Guo. (filed in USPTO, on behalf of VMware). (<https://patents.google.com/patent/US20170371777A1/en>)
- US 2018/0088853 A1. "H/W-S/W co-design for heterogeneous memory management". **Jagadish Kotra**, Alaa A., Chris Wilkerson, Jaewoong S. (<https://patents.google.com/patent/US20180088853A1/en>), Intel.
- US8627230B2. "Intelligent Command Prediction". **Jagadish Kotra**, Anuja Deedwaniya, Shayne Grant, et al. Granted-2014. (<https://patents.google.com/patent/US8627230>), IBM.

TEACHING EXPERIENCE

- Served as a Teaching Assistant for undergraduate Operating Systems (FALL-2011), graduate Operating Systems (Spring-2011) and beginners programming languages course (FALL-2010)
- Guest Lectured undergrad computer architecture course (FALL-2015).

PROFESSIONAL EXPERIENCE

Researcher, Member of Technical Staff, AMD Research, Austin. [April 2018 – Present]

- Worked on Exa-scale Path Forward Research program funded by DoE. Worked on projects involving GPU virtual memory optimizations and CPU critical load prediction.
- Several patent applications were filed in USPTO by AMD and several papers are under submission.

Post-doctoral Researcher, AMD Research, Austin. [Sept 2017 – March 2018]

- Worked on Exa-scale Path Forward project funded by DoE. Work included characterizing MPI applications for identifying hardware bottlenecks.
- Worked on optimizing CPU front-end features to improve performance and energy efficiency. Paper published in MICRO-2020.

Research Intern, Intel Labs, Oregon. Mentors: Alaa R. Alameldeen, Zeshan C. [Jan 2016 – May 2016]

- Worked on a hardware-software co-design for heterogeneous memory management. (Patent filed on fast-track in USPTO. Paper accepted in MICRO 2018.
- Worked on a hardware-software co-design to optimize DRAM refresh overheads. Paper accepted in ASPLOS 2017.

Performance Intern, VMware Performance Team, CA. Mentor: S. Kim. [June 2015 – Sept 2015]

- Worked on congestion-aware memory management in VMware ESXi. Proposed a novel dynamic latency probing algorithm that detects congestion in a NUMA system. Proposed and evaluated congestion-

aware memory allocation/migration techniques in commercial VMware ESXi hypervisor. This work is currently part of commercial VMware ESXi. Patent filed in USPTO, paper accepted in IISWC 2017.

- Graduate Intern, Intel Micro-server Team, Oregon.** Mentor: Brinda Ganesh. [May 2013 – Aug 2013]
- Worked on System Agent (SA) of Intel's future SoC product. SA is a common interface to main memory for the requests coming from Core and I/O devices. I implemented an out-of-order processing for requests going to the main memory at the System Agent to address head-of-queue blocking to better utilize memory resources. This out-of-order processing led to a 17% increase in the performance for the micro-server evaluated using a simulator.

- Systems Software Engineer, IBM Software Labs, India.** [June 2006 – July 2010]
- Worked as IBM JVM developer. As part of the JVM team, I worked on development of various JVM components like Java Class loading. I also worked on IBM middleware products like IBM WMQ and Message Broker on various OS's like z/OS (Mainframes) and various UNIX flavors.

HONORS AND AWARDS (SELECTIVELY LISTED)

- Best Paper Nomination, MICRO-2016.
- Assisted in writing 2 NSF proposals at Penn State.
- IBM Invention Achievement Awards, IBM Labs.
- IBM Thank-you awards, IBM Labs

PROFESSIONAL SERVICE

- Student Research Competition (SRC) Chair: CGO-2019.
- Technical Program Committee Member: MASCOTS-2018, ICCD-2019, ICPP-2019, HPCA-2020 (Industry), MICRO-2020.
- External Review Committee Member: ASPLOS-2019/2020, HPCA-2019, ISCA-2020.
- Fundraising committee Member: MICRO-2019. Web-chair for AIM Workshop – 2017.
- Reviewer of TPDS, TCAD, TACO, TC, TODEAS Journals.

STUDENTS MENTORED

Soheil Khadirsharbiyani , PhD Student, Penn State.	[Jan 2019 – Present]
Armin Vakil , PhD Student, Penn State.	[March 2019 – Present]
Minli Liao , PhD Student, Penn State.	[July 2020 – Present]
Anup Sarma , PhD Student, Penn State.	[2018]